

- [7] M. Faruqui, Y. Tsvetkov, D. Yogatama, C. Dyer, and N. Smith, "Sparse overcomplete word vector representations," 2015. [Online]. Available: <https://arxiv.org/abs/1506.02004>
- [8] G. Goh, "Decoding the thought vector," <https://gabgoh.github.io/ThoughtVectors/>, 2016.
- [9] S. Arora, Y. Li, Y. Liang, T. Ma, and A. Risteski, "Linear algebraic structure of word senses, with applications to polysemy," 2018. [Online]. Available: <https://arxiv.org/abs/1601.03764>
- [10] A. Zou, L. Phan, S. Chen, J. Campbell, P. Guo, R. Ren, A. Pan, X. Yin, M. Mazeika, A.-K. Dombrowski, S. Goel, N. Li, M. J. Byun, Z. Wang, A. Mallen, S. Basart, S. Koyejo, D. Song, M. Fredrikson, J. Z. Kolter, and D. Hendrycks, "Representation engineering: A top-down approach to ai transparency," 2023. [Online]. Available: <https://arxiv.org/abs/2310.01405>
- [11] A. M. Turner, L. Thiergart, G. Leech, D. Udel, J. J. Vazquez, U. Mini, and M. MacDiarmid, "Steering language models with activation engineering," 2024. [Online]. Available: <https://arxiv.org/abs/2308.10248>
- [12] N. Panickssery, N. Gabrieli, J. Schulz, M. Tong, E. Hubinger, and A. M. Turner, "Steering llama 2 via contrastive activation addition," 2024. [Online]. Available: <https://arxiv.org/abs/2312.06681>
- [13] A. Ardit, O. Obeso, A. Syed, D. Paleka, N. Panickssery, W. Gurnee, and N. Nanda, "Refusal in language models is mediated by a single direction," 2024. [Online]. Available: <https://arxiv.org/abs/2406.11717>
- [14] N. Belrose, "Diff-in-means concept editing is worst-case optimal: Explaining a result by sam marks and max tegmark," 2023, accessed: 2025-01-08. [Online]. Available: <https://blog.eleuther.ai/diff-in-means/>
- [15] J. Engels, E. J. Michaud, I. Liao, W. Gurnee, and M. Tegmark, "Not all language model features are linear," 2024. [Online]. Available: <https://arxiv.org/abs/2405.14860>
- [16] R. Taori, I. Gulrajani, T. Zhang, Y. Dubois, X. Li, C. Guestrin, P. Liang, and T. B. Hashimoto, "Stanford alpaca: An instruction-following llama model," https://github.com/tatsu-lab/stanford_alpaca, 2023. [Online]. Available: https://github.com/tatsu-lab/stanford_alpaca
- [17] M. Labonne, "harmless_alpaca," https://huggingface.co/datasets/mlabonne/harmless_alpaca, 2024. [Online]. Available: https://huggingface.co/datasets/mlabonne/harmless_alpaca
- [18] A. Zou, Z. Wang, J. Z. Kolter, and M. Fredrikson, "Universal and transferable adversarial attacks on aligned language models," 2023.
- [19] M. Labonne, "harmful_behaviors," https://huggingface.co/datasets/mlabonne/harmful_behaviors, 2024. [Online]. Available: https://huggingface.co/datasets/mlabonne/harmful_behaviors
- [20] Meta, "Model cards and prompt formats - llama 3.2," https://www.llama.com/docs/model-cards-and-prompt-formats/llama3_2, 2023.
- [21] B. W. et al., "Bloom: A 176b-parameter open-access multilingual language model," 2023. [Online]. Available: <https://arxiv.org/abs/2211.05100>
- [22] A. Y. et al., "Qwen2 technical report," 2024. [Online]. Available: <https://arxiv.org/abs/2407.10671>
- [23] M. Labonne, "Uncensor any llm with ablation," *Hugging Face Community Blog*, June 2024, published June 13, 2024. [Online]. Available: <https://huggingface.co/blog/mlabonne/ablation>
- [24] N. Nanda and J. Bloom, "Transformerlens," <https://github.com/TransformerLensOrg/TransformerLens>, 2022.
- [25] K. Pearson, "On lines and planes of closest fit to systems of points in space," *Philosophical Magazine*, vol. 2, no. 11, pp. 559–572, 1901.
- [26] L. van der Maaten and G. Hinton, "Visualizing data using t-SNE," *Journal of Machine Learning Research*, vol. 9, pp. 2579–2605, 2008.
- [27] L. McInnes, J. Healy, and J. Melville, "Umap: Uniform manifold approximation and projection for dimension reduction," *arXiv preprint arXiv:1802.03426*, 2018.
- [28] A. Schilling, A. Maier, R. Gerum, C. Metzner, and P. Krauss, "Quantifying the separability of data classes in neural networks," *Neural Networks*, vol. 139, pp. 278–293, July 2021, epub 2021 Apr 5.
- [29] N. M. Kirch, S. Field, and S. Casper, "What features in prompts jailbreak llms? investigating the mechanisms behind attacks," 2024. [Online]. Available: <https://arxiv.org/abs/2411.03343>
- [30] R. Greenblatt, C. Denison, B. Wright, F. Roger, M. MacDiarmid, S. Marks, J. Treutlein, T. Belonax, J. Chen, D. Duvenaud, A. Khan, J. Michael, S. Mindermann, E. Perez, L. Petrini, J. Uesato, J. Kaplan, B. Shlegeris, S. R. Bowman, and E. Hubinger, "Alignment faking in large language models," 2024. [Online]. Available: <https://arxiv.org/abs/2412.14093>
- [31] F. Hildebrandt, "Refusal-llms," 2025, accessed: 2025-01-11. [Online]. Available: <https://github.com/FabianHildebrandt/Refusal-LLMs>